

Computing Large Scale Linear Least Squares with SKiLLS

Zhen Shao

December 9, 2020

1 Overview

SKiLLS (SKetchIng-Linear-Least-Squares) is a C++ package for finding solutions to over-determined linear least square problems. SKiLLS uses a modern dimensionality reduction technique called sketching, and is particularly suited for large scale linear least squares where the number of measurements/observations is far greater than the number of variables.

Mathematically, SKiLLS solves

$$\min_{x \in \mathbb{R}^d} f(x) := \|Ax - b\|_2^2, \quad (1)$$

where $A \in \mathbb{R}^{n \times d}$ and $b \in \mathbb{R}^n$ given. The matrix A is allowed to be rank-deficient or nearly rank-deficient.

1.1 When to use SKiLLS

Problem (1) frequently appears as a sub-problem in many computational and data science problems, for example, in many non-linear/general function minimisation routine linearises the problem and iteratively solve problems of the form (1). In the following situations using SKiLLS yields significant computational advantage comparing to state-of-the-arts¹:

1. A is dense, large, sufficiently over-determined and (approximately) rank-deficient or with unknown rank.
2. A is dense, large, sufficiently over-determined with high coherence².
3. A is sparse, large, sufficiently over-determined with non-zero entries randomly distributed.
4. A is moderately sparse (e.g. one percent of entries are non-zero), large and moderately over-determined.

In the following situations using SKiLLS is competitive with state-of-the-arts

1. A is dense, large, sufficiently over-determined with full rank.
2. A is sparse, large, moderately over-determined.

¹Blendenpik [1] for dense problems, SPQR [2] and Cholesky-preconditioned LSQR [5] for sparse problems

²The coherence of a matrix A is defined as the largest Euclidean norm of U , where U is defined in the reduced singular value decomposition of A .

Table 1 shows that SKiLLS is more robust than Blendenpik when the matrix A is approximately rank-deficient. Figure 1 shows that SKiLLS outperforms Blendenpik when the matrix A is (full rank) with high coherence. Figure 2 and Figure 3 shows that SKiLLS is competitive with Blendenpik when the matrix A is (full rank).

Figure 4 shows that SKiLLS runs more than 10 times faster than current state-of-the-art sparse solvers on random sparse matrices. Figure 8 shows that SKiLLS significantly outperforms its competitors when A is moderately sparse, large and moderately over-determined. Figure 5, Figure 7 and Figure 6 showcase the competitiveness of SKiLLS on sparse, large and moderately over determined problems.

For further details of computational advantages, see our paper. *Computing Large Scale Linear Least Squares with SKiLLS*, C. Cartis and Z. Shao.

2 Installing SKiLLS

2.1 Dependency

2.2 Installation instruction

2.3 Test the installation

3 Using SKiLLS

This is a library of two/three routines etc....

3.1 Data structure

Maybe point out to the header file (or quote the header file).

Dense Vector SKiLLS uses C++ class *Vec* for dense vectors. To create a $n \times 1$ vector from already existed data, take a pointer to the data and call *Vec*(n , **data*).

Dense Matrix SKiLLS uses C++ class *Mat* for dense matrices. To create a $n \times d$ matrix from already existed data, take a pointer to the data and call *Mat*(n , d , **data*). Basic matrix operations such as return a specific row, column, or matrix-matrix and matrix vector multiplications are implemented.

Sparse Matrix SKiLLS uses Compressed Column Data structure for sparse matrices and vectors. The sparse data structure is from SuiteSparse [3]. In particular, the *cs_dl* structure used as the input matrix format for sparse linear least squares is taken from the CXSparse package from SuiteSparse. It is a C structure with 7 fields:

- *nzmax*: maximum number of non-zeros.
- *m*: number of rows.
- *n*: number of columns.
- **p*: column indices (size *nzmax*)
- **i*: row indices, size *nzmax*.
- **x*: numerical values of non-zeros, size *nzmax*.
- *nz*: -1 for compressed column format.

Sometimes we will also use the data structure `cholmod_sparse` for sparse matrix in compressed column format. It has a similar structures. For more information, see the header file or documentation of SuiteSparse.

Examples We give two examples of constructing dense and sparse matrices from user given data. The below code snippet creates a dense matrix and a vector.

```
long m = 4;
long n = 2;
double data_A[8] = {0.306051, 1.53078, 1.64493, 1.61322,
                    0.2829, 0.474476, 0.586278, 0.610202};

double data_b[4] = {0.0649338, 0.845946, 0.0164085, 0.247119};

Mat_d A(m,n,data_A);
Vec_d b(m,data_b);
```

The below code snippet creates a sparse matrix.

```
long m = 4;
long n = 3;
long nnz = 5;

long col[4] = {1, 2, 4, 6};
long row[5] = {1, 1, 3, 2, 4};
double val[5] = {2.0, 3.0, 1.0, 4.0, 5.0};
long col_0[4];
long row_0[5];

// convert to zero based indices
for (long i=0; i<n+1; i++){
    col_0[i] = col[i] - 1;
}

for (long i =0; i<nnz; i++){
    row_0[i] = row[i] - 1;
}

cholmod_common Common, *cc;
cc = &Common;
cholmod_l_start(cc);
cc->print = 4;

cholmod_sparse* A;
A = cholmod_l_allocate_sparse(
m, n, nnz, true, true, 0, CHOLMOD_REAL, cc);

A->p = col_0;
A->i = row_0;
A->x = val;

cholmod_l_print_sparse( A, "A", cc);

cholmod_l_finish(cc);
```

3.2 Dense problems

To compute a solution of a linear least square problem where the matrix A is dense, a call of the following form should be made.

`ls_dense_hashing_blendenpik(A, b, x, rank, flag, it, gamma, k, abs_tol, rcond, it_tol, max_it, debug, wisdom):`

separate into input, output and options typical correct values of these parameters

- A is an input of derived type containing the $\mathbb{R}^{m \times n}$ matrix.
- b is an input of derived type containing the \mathbb{R}^m vector.
- x is an output of derived type containing the \mathbb{R}^n vector solution on exit.
- $rank$ is an output of derived type containing a scalar, detected rank of the matrix A .
- $flag$ indicates LSQR convergence.
- it indicates LSQR iteration count.
- $gamma$ is the over-sampling ratio used in sketching.
- k is number of non-zeros per column in the hashing matrix.
- abs_tol is absolute residual tolerance, the algorithm will immediately return if it finds a vector x with $\|Ax - b\| \leq abs_tol$.
- $rcond$ a parameter used in column pivoted QR to determine rank. After a column pivoted QR factorization of the matrix $SAP = QR$, diagonal entries of R less than $rcond$ is treated as zero and the whole corresponding column in the matrix Q is discarded.
- it_tol is tolerance used in the preconditioned LSQR algorithm for LSQR convergence.
- max_it is maximum iteration allowed in the preconditioned LSQR algorithm.
- $debug$ is a flag. If $debug=1$ then the solver will print some additional information.
- $wisdom$ is a flag. If $wisdom=1$ then the user is expected to have provided a wisdom file for the Discrete Cosine Transform used in sketching.

If the matrix A is known to be of full numerical rank, then a different routine can be used which is slightly faster:

`ls_dense_hashing_blendenpik_noCPQR(A, b, x, rank, flag, it, gamma, k, it_tol, max_it, debug, wisdom):`

The arguments usage is the same as above, but we don't need the rank-detection parameter $rcond$, and the shortcut related to abs_tol is not implemented.

Example The following program solves an example of problem (1) with given A and b .

The matrix A is created by

```
#include <iostream>
#include "Config.hpp"
#include "SpMat.hpp"
#include "Random.hpp"
#include <math.h>
#include "RandSolver.hpp"
#include "cs.h"
#include <stdio.h>
#include "SuiteSparseQR.hpp"
#include "cholmod.h"
#include "lsex_all_in_c.h"
#include "IterSolver.hpp"
#include <sys/timeb.h>
#include "bench_config.hpp"
#include "blendenpik.hpp"
```

```

int main(int argc, char **argv)
{
    // create data
    long m = 4;
    long n = 2;
    double data_A[8] = {0.306051, 1.53078, 1.64493, 1.61322,
        0.2829, 0.474476, 0.586278, 0.610202};

    double data_b[4] = {0.0649338, 0.845946, 0.0164085, 0.247119};

    Mat_d A(m, n, data_A);
    Vec_d b(m, data_b);

    // parameters
    long k = NNZ_PER_COLUMN;
    double gamma = OVER_SAMPLING_RATIO;
    long max_it = MAX_IT;
    double it_tol = IT_TOL;
    double rcond = 1e10;
    int wisdom = 0;
    double abs_tol = ABS_TOL;

    // storage
    long it;
    int flag;
    long rank;
    double residual;
    Vec_d x(A.n());

    // solve
    ls_dense_hashing_blendenpik(
        A, b, x, rank, flag, it, gamma, k, abs_tol,
        rcond, it_tol, max_it, debug, wisdom);

    std::cout << "A_is:_ "<< A << std::endl;
    std::cout << "b_is:_ "<< b << std::endl;
    std::cout << "solution_is:_ "<< x << std::endl;

    // compute residual
    A.mv('n', 1, x, 1, b);
    std::cout << "Residual_of_ls_blendenpik_hashing_is:_ "<< nrm2(b) << std::endl;
    std::cout << "Iteration_of_ls_blendenpik_hashing_is:_ " << it << std::endl;
}

```

3.3 Sparse problems

To compute a solution of a linear least square problem where the matrix A is sparse, a call of the following form should be made.

`ls_sparse_spqr(*A, b, x_ls_qr, rank, flag, it, gamma, k, abs_tol, ordering, it_tol, max_it, rcond, peturb, debug);`

- A is an input of derived type containing the $\mathbb{R}^{m \times n}$ matrix.
- b is an input of derived type containing the \mathbb{R}^m vector.

- *x* is an output of derived type containing the \mathbb{R}^n vector solution on exit.
- *rank* is an output of derived type containing a scalar, detected rank of the matrix *A*.
- *flag* indicates LSQR convergence.
- *it* indicates LSQR iteration count.
- *gamma* is the over-sampling ratio used in sketching.
- *k* is number of non-zeros per column in the hashing matrix.
- *abs_tol* is absolute residual tolerance, the algorithm will immediately return if it finds a vector *x* with $\|Ax - b\| \leq \text{abs_tol}$.
- *ordering* is a parameter used in to select fill-reduced ordering strategy in sparse QR factorization.
- *it_tol* is tolerance used in the preconditioned LSQR algorithm for LSQR convergence.
- *max_it* is maximum iteration allowed in the preconditioned LSQR algorithm.
- *rcond* is a parameter used in sparse rank-revealing QR to determine rank.
- *perturb* is a potential pivot perturbation if R_{11} returned by the sparse QR factorization is ill-conditioned.
- *debug* is a flag. If *debug*=1 then the solver will print some additional information.

4 Reference

References

- [1] H. Avron, P. Maymounkov, and S. Toledo. *Blendenpik: Supercharging LAPACK’s Least-Squares Solver*. SIAM Journal on Scientific Computing, 32(3):1217–1236, Jan. 2010.
- [2] T. A. Davis. *Algorithm 915, SuiteSparseQR: Multifrontal multithreaded rank-revealing sparse QR factorization*. ACM Transactions on Mathematical Software (TOMS), 38(1):8, Nov. 2011.
- [3] T. A. Davis. *Algorithm 915, SuiteSparseQR: Multifrontal Multithreaded Rank-Revealing Sparse QR Factorization*. ACM Transactions on Mathematical Software (TOMS), 38(1), Dec. 2011.
- [4] T. A. Davis and Y. Hu. *The University of Florida Sparse Matrix Collection*. ACM Transactions on Mathematical Software (TOMS), 38(1), Dec. 2011.
- [5] J. Scott and M. Tũma. *HSL_MI28: An Efficient and Robust Limited-Memory Incomplete Cholesky Factorization Code*. ACM Transactions on Mathematical Software, 40(4):1–19, June 2014.

A Numerical Results

	lp_ship12l	Franz1	GL7d26	cis-n4c6-b2	lp_modszk1	rel5	ch5-5-b1
SVD	18.336	26.503	50.875	6.1E-14	33.236	14.020	7.3194
Blendenpk	NaN	9730.700	NaN	3.0E+02	NaN	NaN	340.9200
Ski-LLS-dense	18.336	26.503	50.875	5.3E-14	33.236	14.020	7.3194
	n3c5-b2	ch4-4-b1	n3c5-b1	n3c4-b1	connectus	landmark	cis-n4c6-b3
SVD	9.0E-15	4.2328	3.4641	1.8257	282.67	1.1E-05	30.996
Blendenpk	1.3E+02	66.9330	409.8000	8.9443	NaN	NaN	3756.200
Ski-LLS-dense	5.2E-15	4.2328	3.4641	1.8257	282.67	1.1E-05	30.996

Table 1: Residual of solvers for a range of approximate rank-deficient problems taken from the Florida matrix collection [4]. We see while Blendenpik struggles, SKiLLS achieves the same residual accuracy as SVD method.




Figure 1: Time taken by solvers to compute the solution of problem (1) for A being (full rank) coherent dense matrices of various sizes (x-axis)




Figure 2: Time taken by solvers to compute the solution of problem (1) for A being (full rank) semi-coherent dense matrices of various sizes (x-axis)




Figure 3: Time taken by solvers to compute the solution of problem (1) for A being (full rank) incoherent dense matrices of various sizes (x-axis)

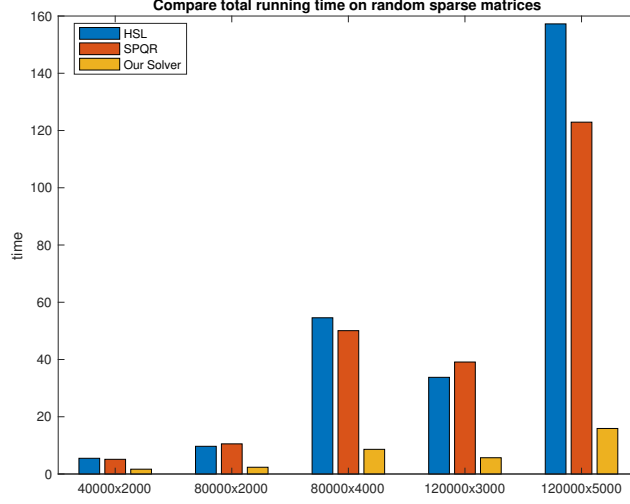


Figure 4: Our solver is much faster than state-of-the-arts LS_HSL and LS_SPQR on a range of random sparse data matrices.




Figure 5: Performance profile comparison of Ski-LLS with LSRN, LS_HSL and LS_SPQR for all matrices $A \in \mathbb{R}^{n \times d}$ in the Florida matrix collection with $n \geq 30d$.




Figure 6: Performance profile comparison of Ski-LLS with LSRN, LS_HSL and LS_SPQR for all matrices $A \in \mathbb{R}^{n \times d}$ in the Florida matrix collection with $n \geq 10d$.




Figure 7: Performance profile comparison of Ski-LLS with LSRN, LS_HSL and LS_SPQR for all matrices $A \in \mathbb{R}^{n \times d}$ in the Florida matrix collection with $n \geq 10d$ and the unpreconditioned LSQR takes more than 5 seconds to solve.




Figure 8: All solvers, aspect ratio at least 10 and more than 1 percent number of non-zeros